

Exploring Census Data

Emma Johnson

December 2025

S1201: Marital Status

```
browseURL("https://data.census.gov/table/ACSST1Y2024.S1201?g=010XX00US$0400000")
library(tidyverse)

## -- Attaching core tidyverse packages ----- tidyverse 2.0.0 --
## v dplyr     1.1.4     v readr     2.1.5
## vforcats   1.0.0     v stringr   1.5.1
## v ggplot2   3.5.2     v tibble    3.3.0
## v lubridate 1.9.4     v tidyrr    1.3.1
## v purrr    1.1.0

## -- Conflicts ----- tidyverse_conflicts() --
## x dplyr::filter() masks stats::filter()
## x dplyr::lag()   masks stats::lag()
## i Use the conflicted package (<http://conflicted.r-lib.org/>) to force all conflicts to become errors

library(tidycensus)

## Warning: package 'tidycensus' was built under R version 4.5.2
```

Here is the typical set-up to fetch census data from the United States Census Bureau:

```
#census_api_key(API_key, install = TRUE)
## Use your API key ^ here

marital_status <- get_acs(geography = "state", table = "S1201")

## Getting data from the 2019-2023 5-year ACS

## Loading ACS5/SUBJECT variables for 2023 from table S1201. To cache this dataset for faster access to

## Using the ACS Subject Tables
```

Notice how the variables are not informative under their current codes, so we will now fetch more descriptive names using `load_variables()`. To keep the data as concise as possible, we will use these variable names in our main dataframe, `ms` (marital status).

```
ms <- marital_status %>%
  left_join(
    load_variables(2024, "acs1/subject"),
    join_by(variable == name),
    keep = FALSE,
    relationship = "many-to-one"
  ) %>%
  select(-concept)

glimpse(ms)
```

By taking a glimpse of the data, you may notice that the variable names are long and inefficiently named. This is because the data has levels of headers, which can make calling a variable quite difficult. So, the easiest way I found to fix this was to ask ChatGPT to clean it up. Otherwise, feel free to have at it yourself.

```
## With a little clean-up help from ChatGPT:  
ms$clean <- ms$label %>%  
  sub("^Estimate!!", "", .) %>%  
  str_split("!!") %>%  
  sapply(function(parts) {  
    parts <- trimws(tolower(parts))  
  
    # ---- 1. marital status ----  
    status <- parts[1]  
  
    # normalize "now married (except separated)" to "married"  
    status <- gsub("^now married.*", "married", status)  
  
    # ---- 2. last meaningful piece ----
```

```

last <- parts[length(parts)]

# ---- 3. detect sex (male/female) anywhere ----
sex_word <- NA_character_
if (any(grepl("\\\\bmales?\\\\b", parts))) sex_word <- "male"
if (any(grepl("\\\\bfemales?\\\\b", parts))) sex_word <- "female"

# ---- 4. remove male/female from final chunk to avoid duplicates ----
if (!is.na(sex_word)) {
  last <- gsub("\\\\bmales?\\\\b", "", last)
  last <- gsub("\\\\bfemales?\\\\b", "", last)
  label <- paste(status, sex_word, last)
} else {
  label <- paste(status, last)
}

# ---- 5. snake_case cleanup ----
label <- gsub("[^a-z0-9]+", "_", label)
label <- gsub("_+", "_", label)
label <- gsub("^_|_$", "", label)

label
})

ms <- ms %>%
  select(geo_id = GEOID, name = NAME, variable = clean, estimate, error_margin = moe)

glimpse(ms)

```

```

## Rows: 9,984
## Columns: 5
## $ geo_id      <chr> "01", "01", "01", "01", "01", "01", "01", "01", "01~
## $ name        <chr> "Alabama", "Alabama", "Alabama", "Alabama", "A~
## $ variable    <chr> "total_population_15_years_and_over", "total_male_15_year~
## $ estimate     <dbl> 4124301, 1978662, 172319, 488248, 303517, 303401, 321761, ~
## $ error_margin <dbl> 773, 1057, 1324, 1207, 985, 1050, 569, 474, 873, 1196, 13~

```

There is lots of redundancy in the data due to grouping by state. That is, columns `geo_id` and `name` (i.e. state/territory name) are repeated for each of their corresponding observations. To improve efficiency, we will want to use this grouping to our advantage, rather than having to make the system re-group every time we want to look at specific states. So, we will nest each territory's data into it's own dataframe using `group_nest` and `group_keys` for lookup.

```

ms_by_state <- ms %>%
  group_by(geo_id, name) %>%
  group_nest(keep = FALSE)

state_keys <- ms %>%
  group_by(geo_id, name) %>%
  group_keys()

# e.g. Alabama
ms_by_state$data[[1]]

```

```

## # A tibble: 192 x 3
##   variable                estimate error_margin
##   <chr>                  <dbl>      <dbl>
## 1 total_population_15_years_and_over 4124301     773
## 2 total_male_15_years_and_over       1978662    1057
## 3 total_male_15_to_19_years         172319      1324
## 4 total_male_20_to_34_years         488248      1207
## 5 total_male_35_to_44_years         303517      985
## 6 total_male_45_to_54_years         303401      1050
## 7 total_male_55_to_64_years         321761      569
## 8 total_male_65_years_and_over     389416      474
## 9 total_female_15_years_and_over   2145639     873
## 10 total_female_15_to_19_years      169532      1196
## # i 182 more rows

```

It's looking much more digestible already. In addition to the grouping by state, there is also latent groupings of variable type. To optimize efficiency once more, I created an index for the types of variables, so finding them will be easier later.

```

variable_index <- list(
  total_pop = ms$variable[1:15],
  race = ms$variable[16:26],
  labor_force = ms$variable[27:30],
  ratio_sex = ms$variable[31],

  marital_status_percentage = ms$variable[32],
  married_pop = ms$variable[33:47],
  married_race = ms$variable[48:58],
  married_work_force = ms$variable[59:62],

  widowed_pop = ms$variable[65:79],
  widowed_race = ms$variable[80:90],
  widowed_work_force = ms$variable[91:94],

  divorced_pop = ms$variable[97:111],
  divorced_race = ms$variable[112:122],
  divorced_work_force = ms$variable[123:126],

  separated_pop = ms$variable[129:143],
  separated_race = ms$variable[144:154],
  separated_work_force = ms$variable[155:158],

  never_married_pop = ms$variable[161:175],
  never_married_race = ms$variable[176:186],
  never_married_work_force = ms$variable[187:190]
)

```

As an example of what we've cleaned up so far, here is how we can fetch a group of variables for a specific state, or even for all states in a comparison:

```

# Widowed population variables in Alabama (2 ways):
ms_by_state$data[[1]] %>%
  filter(variable %in% variable_index$widowed_pop)

```

```

## # A tibble: 16 x 3
##   variable           estimate error_margin
##   <chr>              <dbl>      <dbl>
## 1 widowed_population_15_years_and_over    6.9       0.1
## 2 widowed_male_15_years_and_over          3.3       0.1
## 3 widowed_male_15_to_19_years            0         0.1
## 4 widowed_male_20_to_34_years           0.1       0.1
## 5 widowed_male_35_to_44_years           0.4       0.1
## 6 widowed_male_45_to_54_years           1.5       0.2
## 7 widowed_male_55_to_64_years           3.5       0.3
## 8 widowed_male_65_years_and_over        12.2      0.4
## 9 widowed_female_15_years_and_over      10.2      0.2
## 10 widowed_female_15_to_19_years        0         0.1
## 11 widowed_female_20_to_34_years        0.3       0.1
## 12 widowed_female_35_to_44_years        1.2       0.2
## 13 widowed_female_45_to_54_years        3.4       0.4
## 14 widowed_female_55_to_64_years        9.3       0.4
## 15 widowed_female_65_years_and_over     34.5      0.5
## 16 widowed_population_15_years_and_over  6.9       0.1

```

```

# OR
ms_by_state %>%
  filter(name == "Alabama") %>%
  unnest(cols = everything()) %>%
  filter(variable %in% variable_index$widowed_pop)

```

```

## # A tibble: 16 x 5
##   geo_id name   variable           estimate error_margin
##   <chr>  <chr>  <chr>              <dbl>      <dbl>
## 1 01     Alabama widowed_population_15_years_and_over    6.9       0.1
## 2 01     Alabama widowed_male_15_years_and_over          3.3       0.1
## 3 01     Alabama widowed_male_15_to_19_years            0         0.1
## 4 01     Alabama widowed_male_20_to_34_years           0.1       0.1
## 5 01     Alabama widowed_male_35_to_44_years           0.4       0.1
## 6 01     Alabama widowed_male_45_to_54_years           1.5       0.2
## 7 01     Alabama widowed_male_55_to_64_years           3.5       0.3
## 8 01     Alabama widowed_male_65_years_and_over        12.2      0.4
## 9 01     Alabama widowed_female_15_years_and_over      10.2      0.2
## 10 01    Alabama widowed_female_15_to_19_years        0         0.1
## 11 01    Alabama widowed_female_20_to_34_years        0.3       0.1
## 12 01    Alabama widowed_female_35_to_44_years        1.2       0.2
## 13 01    Alabama widowed_female_45_to_54_years        3.4       0.4
## 14 01    Alabama widowed_female_55_to_64_years        9.3       0.4
## 15 01    Alabama widowed_female_65_years_and_over     34.5      0.5
## 16 01    Alabama widowed_population_15_years_and_over  6.9       0.1

```

```

# Ratio of Unmarried Men 15 to 44 years per 100 unmarried women 15 to 44 years
ms %>%
  filter(variable %in% variable_index$ratio_sex) %>%
  select(name, estimate, error_margin)

```

```

## # A tibble: 52 x 3
##   name           estimate error_margin
##   <chr>          <dbl>      <dbl>
## 1
## 2
## 3
## 4
## 5
## 6
## 7
## 8
## 9
## 10
## 11
## 12
## 13
## 14
## 15
## 16
## 17
## 18
## 19
## 20
## 21
## 22
## 23
## 24
## 25
## 26
## 27
## 28
## 29
## 30
## 31
## 32
## 33
## 34
## 35
## 36
## 37
## 38
## 39
## 40
## 41
## 42
## 43
## 44
## 45
## 46
## 47
## 48
## 49
## 50
## 51
## 52

```

```
##      <chr>          <dbl>        <dbl>
## 1 Alabama       104.         0.7
## 2 Alaska        131.         4.2
## 3 Arizona       115.         0.6
## 4 Arkansas      110.         1
## 5 California    113.         0.2
## 6 Colorado       120.         0.7
## 7 Connecticut   108.         0.7
## 8 Delaware      102.         1.4
## 9 District of Columbia 87.3     1
## 10 Florida      110.         0.4
## # i 42 more rows
```